

# Linear models

---

**Alessio Micheli**

**[micheli@di.unipi.it](mailto:micheli@di.unipi.it)**



Dipartimento di Informatica  
Università di Pisa - Italy

**Computational Intelligence &  
Machine Learning Group**

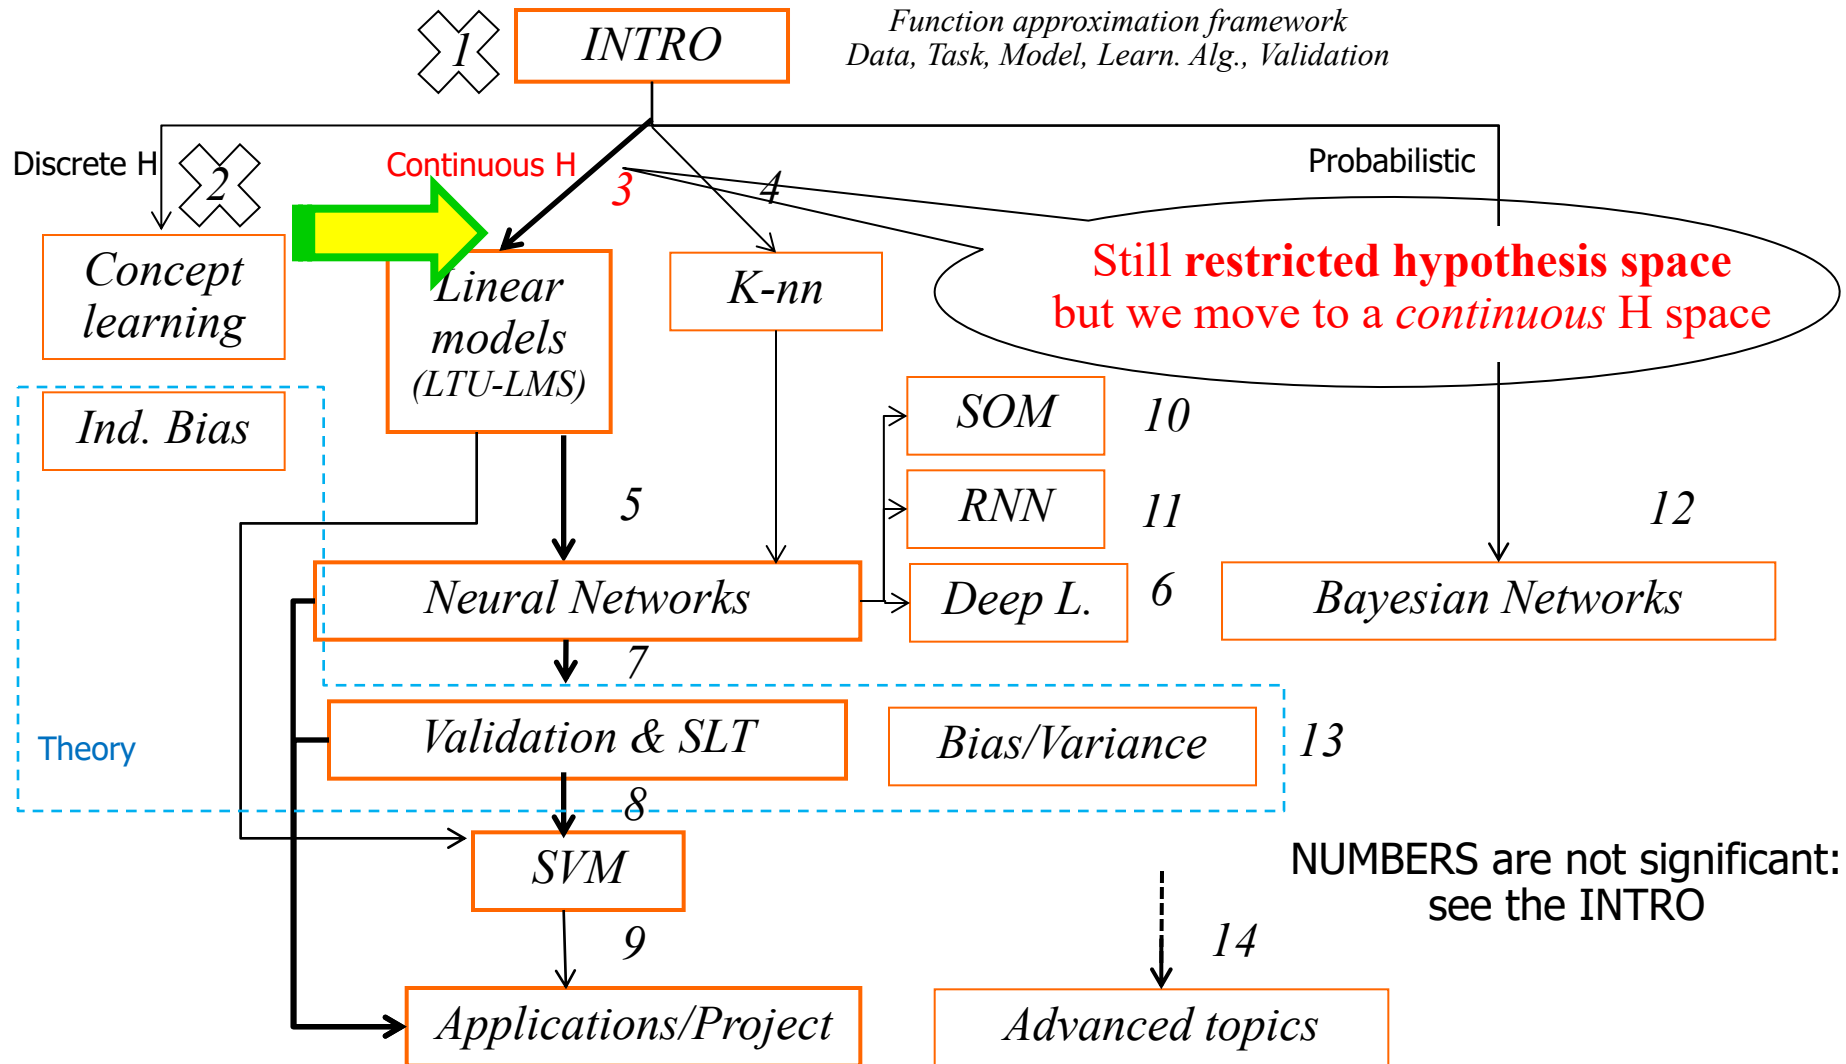
*Sept, 2025*

# ML Course structure

## Where we go




Dip. Informatica  
University of Pisa



# Tasks: Supervised Learning

- **Given:** Training examples as  $\langle input, output \rangle = (x, d)$  (*labeled examples*)  
for an unknown function  $f$  (known only at the given points of example)
  - Target value: desiderate value  $d$  or  $t$  or  $y$  ... is given by the teacher according to  $f(x)$ .
- **Find:** A *good* approximation to  $f$  (a hypothesis  $h$  that can be used for prediction on unseen data  $x'$ , i.e. that is able to generalize)



- Target  $d$  (or  $t$  or  $y$ ): a categorical or numerical *label*
  - **Classification:** discrete value outputs:  
 $f(x) \in \{1, 2, \dots, K\}$  *classes (discrete-valued function)*
  - **Regression:** real continuous output values (approximate a real-valued target function)

**Both as a *task of function approximation***

# A premise on **DATA** notation

Pattern	$x_1$	$x_2$	$x_i$	$x_n$
Pat 1	$x_{1,1}$	$x_{1,2}$		$x_{1,n}$
...				
Pat $p$	$x_{p,1}$	$x_{p,2}$	$x_{p,i}$	$x_{p,n}$
...				

$X$  is a matrix  $l \times n$

$l$  rows,  $n$  columns

$p=1..l, \quad i=1..n$

We often need to omit some indices when the context is clear, e.g.:

- Each row, generic  $x$  (vector - bold), a row in the table: (input) example, pattern, instance, sample, ..., input vector, ...
- $x_i$  or  $x_j$  (scalar): component  $i$  or  $j$  (given a pattern  $x$ , i.e. omitting  $p$ )
- $x_p$  or  $x_i$  (vector – bold)  $p$ -th or  $i$ -th row in the table = pattern  $p$  or  $i$
- $x_{p,i}$  (scalar) also as  $(x_p)_i$ : component  $i$  of the pattern  $p$   
or also, often,  $x_{p,j}$  for the component  $j$ , etc.
- For the target  $y$  we will typically use just  $y_p$  with  $p=1..l$  (the same for  $d$  or  $t$ )

# Linear models

---

- **Regression**
- **Classification**

# Linear models

The linear model has been the mainstay of statistics

- *"Despite the great inroads made by modern nonparametric regression techniques, linear models remain important, and so we need to understand them well". (Hastie)*

Plenty of studies and in many books (mathematics, statistics, numerical analysis, applicative fields, ML, ...)

- We start with the simplest form, linear in the input variables
- A baseline for learning (first: is it a linear problem?)

# Regression

---


- **Just to see how formulate the *learning problem* as a **LMS** on the  $R_{emp}$** 
  - Then (next lecture) we will consider also the control of complexity
- **We formulate a first derivation in a simplified setting (univariate case)**

# Repetita: regression: example

- Process of estimating of a real-value function on the basis of a finite set of noisy samples
  - known pairs  $(x, f(x) + \text{random noise})$

Task (exercise): find  $f$  for the data in the following table:

$x$	target
1	2.1
2	3.9
3	6.1
4	8.4
5	9.8
...	...




Now, we want to solve it (how to find  $w$ ) in a «systematic» way




# Univariate Linear Regression

- Univariate case, simple linear regression :
- We start with 1 input variable  $x$ , 1 output variable  $y$
- We assume a model  $h_w(x)$  expressed as  $out = w_1 x + w_0$
- *where  $w$  are real-valued coefficients/free parameters (weights)*



- **Fitting** the data by a “straight line”



- Infinite hp space (continuous  $w$  values) but we have nice solution from classical math (going back to Gauss/Legendre ~1795!)
  - Surprisingly we can “learn” by this basic tool
  - Although simple it includes many relevant concept of modern ML and it is a basis of evolved methods in the field

# Build it: Learning via LMS (I)

- **Training** → find  $\mathbf{w}$  such that minimize **error**/empirical **loss** (best data fitting – on the training set with  $l$  examples): i.e. we are now focusing on the  $R_{emp}$
- **Given** a set of  $l$  training examples  $(\mathbf{x}_p, y_p)$   $p=1..l$
- **Find:**  $h_{\mathbf{w}}(x)$  in the form  $w_1 x + w_0$  (hence the values of  $\mathbf{w}$ ) that minimizes the expected loss on the training data.
- For the loss we use the square of errors:
- Least (Mean) Square: Find  $\mathbf{w}$  to *minimize* the residual sum of squares  $[argmin_{\mathbf{w}} Error(\mathbf{w}) \text{ in } L_2]$ :

$$Loss(h_{\mathbf{w}}) = E(\mathbf{w}) = \sum_{p=1}^l (y_p - h_{\mathbf{w}}(x_p))^2 = \sum_{p=1}^l (y_p - (w_1 x_p + w_0))^2$$

*p runs over patterns/examples*

where  $x_p$  is  $p$ -th input/pattern/example,  $y_p$  the output for  $p$ ,  $w$  free par.,  $l$  num. of examples


Note: to have the mean divide by  $l$

On the notation: Indeed for the univariate case, with 1 variable:  $x_p = x_{p,1} = (\mathbf{x}_p)_1$

# Build it: Learning via LMS (II)

## Why LMS to fit the data with $h$ ?

- Least (Mean) Square: Find  $w$  to minimize the **residual sum of squares**  
[ $\operatorname{argmin}_w \operatorname{Error}(w)$  in  $L_2$ ]:



$$h_w(x)$$

$$y = w_1 x + w_0 + \text{noise}$$

Different **blue** lines will have different **green** bars.

Minimizing the **green** bars (residuals / errors)

is a way to find the best approximation/fitting of the data (i.e. our  $h_w(x)$  or **blue** line).

The squares of errors  $E(w)$  quantify such **green** bars:

$$E(w) = \sum_{p=1}^l (y_p - h_w(x_p))^2$$

- The method of **least squares** is a standard approach to the approximate solution of over-determined systems, i.e., sets of equations in which there are more equations than unknowns.

# How to solve?


- Remember: local minimum as stationary point: the gradient is zero

$$\frac{\partial E(\mathbf{w})}{\partial w_i} = 0, \quad i = 1, \dots, \dim\_input + 1 = 1, \dots, n + 1$$


- For the simple Lin. Regr. (2 free parameters)

Search the  $\mathbf{w}$  such that

$$\frac{\partial E(\mathbf{w})}{\partial w_0} = 0 \quad \frac{\partial E(\mathbf{w})}{\partial w_1} = 0$$



Convex loss function  $\rightarrow$  we have the following solution (no local minima)  
(just to know that it exists!)



$$w_1 = \frac{\sum_{p:1 \rightarrow l} x_p y_p - \frac{1}{l} \sum x_p \sum y_p}{\sum x_p^2 - \frac{1}{l} (\sum x_p)^2} = \frac{\text{Cov}[x, y]}{\text{Var}[x]},$$

**Exercise:** compute  $w_0$  and  $w_1$  according to the next slide results for the gradient (extended to  $l$  patterns)

$$w_0 = \bar{y} - w_1 \bar{x}$$

$\downarrow$                        $\downarrow$   
 $\frac{1}{l} \sum_{p \rightarrow l} y_p$        $\frac{1}{l} \sum_{p \rightarrow l} x_p$

# Compute the gradient for 1 (each) pattern $p$



Dip. Informatica  
University of Pisa

Redo this by yourself as an *Exercise*

Basic rules:

$$\frac{\partial}{\partial w} k = 0, \frac{\partial}{\partial w} w = 1, \frac{\partial}{\partial w} w^2 = 2w$$

$$\frac{\partial (f(w))^2}{\partial w} = 2f(w) \frac{\partial (f(w))}{\partial w}$$

*Der. sum = sum of der.*

We will call  $(y - h_w(x))$   
"delta"

$$\frac{\partial E(\mathbf{w})}{\partial w_i} = \frac{\partial (y - h_w(x))^2}{\partial w_i} =$$

$$= 2(y - h_w(x)) \frac{\partial (y - h_w(x))}{\partial w_i} = 2(y - h_w(x)) \frac{\partial (y - (w_1 x + w_0))}{\partial w_i}$$

$$\frac{\partial E(\mathbf{w})}{\partial w_0} = -2(y - h_w(x))$$

$$\frac{\partial E(\mathbf{w})}{\partial w_1} = -2(y - h_w(x)) \cdot x$$

Then we will sum up for  $l$  patterns  $(x_p, y_p) \dots$

And we will extend to multidimensional  $x$  and  $w$  ( $n=1$  here) ...

# Linear model: notation for multidimensional inputs

- Assuming column vector for  $\mathbf{x}$  and  $\mathbf{w}$  (in bold)
  - Number of data  $l$ , dimension of input vector  $n$ ,  $y_p$  (targets)  $p=1..l$
- $$\mathbf{w}^T \mathbf{x} + w_0 = w_0 + w_1 x_1 + w_2 x_2 + \dots + w_n x_n = w_0 + \sum_{i=1}^n w_i x_i \quad (\text{eq. 1})$$
- Note that sometimes (in NN) the transpose notation  $T$  in  $\mathbf{w}^T$  is omitted
  - $w_0$  is the *intercept, threshold, bias, offset*....

Often it is convenient to include the constant  $x_0 = 1$  so that we can write eq.1 as :

$$\mathbf{w}^T \mathbf{x} = \boxed{\mathbf{x}^T \mathbf{w}}$$

Inner product

$$\mathbf{x}^T = [1, x_1, x_2, \dots, x_n]$$

$$\mathbf{w}^T = [w_0, w_1, w_2, \dots, w_n]$$

So, the “linear” model can be written  
as a function that for each  $\mathbf{x}_p$  compute:

$$h(\mathbf{x}_p) = \mathbf{x}_p^T \mathbf{w} = \sum_{i=0}^n x_{p,i} w_i$$

$\mathbf{w}$ : continuous (free) parameters:  
“weights”

# Learning algorithm: just wait!



Dip. Informatica  
University of Pisa

- For the learning algorithm in the multidimensional case for linear regression: please wait.
- We will provide it along with the learning algorithm of the linear classifier in the next few slides.

## Summing up:

- Given the data set and the linear model, we can state the learning problem as **LMS problem**
- Once we find the best  $w$  parameters values, we have our  $h_w(x)$  for regression purposes
- For students that need a soft intro to just LMS Regression:
  - <https://svivek.com/teaching/machine-learning/lectures/slides/linear-models/lms-regression.pdf>

# Classification

---




# What we are looking at

---

1. The classification by **hyperplanes**
2. See how the model works after training (“**use it**” slides)
3. How to state/formulate a **(regression)/classification learning problem** for a linear model by LMS
4. How to **derive the learning algorithm**
5. Proposing **two learning algorithms** to build a linear classifier

# Problem: example

Raw Data with a Binary Response



200 points generated in  $\mathbb{R}^2$  from an unknown distribution; 100 in each of two classes.

Can we build a rule to predict the color of future points?

Data may be generated by gaussian distribution (for each class) with different means  
or by a mixture of different low variance gaussian distributions.

# We reuse the linear model


- The same models (used for regression) can be used for **classification**: categorical targets ( $y$  or  $d$ ), e.g. **0/1** or **-1/+1**.
- In this case we use a **hyperplane** ( $w^T x$ ) assuming negative or positive values
- We exploit such models to decide if a point  $x$  belong to positive or negative zone of the hyperplane (to classify it)
- So we want to set  $w$  (by learning) s.t. we get good classification accuracy

# Geometrical view: **hyperplane**

$w^T x$  define an hyperplane.  
E.g. see the picture for 2 variables

$$x^T w = w^T x = w_0 + w_1 x_1 + w_2 x_2 = 0$$

*This defines the decision boundary*



Can be used to classify:

Examples  $\langle (x_1, x_2), y \rangle$ :

$\langle (1.0, 1.0), 1 \rangle$


$\langle (0.5, 3.0), 1 \rangle$

$\langle (2.0, 2.0), 0 \rangle$

# Geometrical view: classifier

Introducing a threshold function

Region where the  
output of the  
classifier is 1



Using  $\mathbf{x}_p$  and  
including  $w_0$  in  $\mathbf{w}$

In this slide  $\mathbf{w}$  is omitted from  $h_{\mathbf{w}}$  (we use just  $h$ )

Examples:

$\langle (1.0, 1.0), 1 \rangle$

$\langle (0.5, 3.0), 1 \rangle$

$\langle (2.0, 2.0), 0 \rangle$

$[0,1]$  output range

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w}\mathbf{x} + w_0 \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

or

$[-1,+1]$  output range

$$h(\mathbf{x}) = \text{sign}(\mathbf{w}\mathbf{x} + w_0)$$


$$h(\mathbf{x}_p) = \text{sign}(\mathbf{x}_p^T \mathbf{w}) = \text{sign}\left(\sum_{i=0}^n x_{p,i} w_i\right)$$

# Classification by linear decision boundary [repetita]

The classification may be viewed as the allocation of the input space in decision regions (e.g. **0/1**)

Example: linear separator on

2-dim instance space  $\mathbf{x}=(x_1, x_2)$  in  $\mathbb{R}^2$ ,  $f(\mathbf{x})=0/1$  (or  $-1/+1$ )



Point belonging to class 1

Separating (hyper)plane :  $\mathbf{x}$  s.t.

$$\mathbf{w}^T \mathbf{x} + w_0 = w_1 x_1 + w_2 x_2 + w_0 = 0$$

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w}^T \mathbf{x} + w_0 \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

[0,1]  
output range

or

$$h(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x} + w_0)$$

[-1,+1]  
output range

**Linear threshold unit (LTU)**

Indicator functions

How many? (H): set of dichotomies induced by hyperplanes

# Threshold (bias $w_0$ )

Note that, given the bias  $w_0$ , in the LTU

saying  $h(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0 \geq 0$

is equivalent to say  $h(\mathbf{x}) = \mathbf{w}^T \mathbf{x} \geq -w_0$

with  $-w_0$  as the «threshold» value

- The two forms identify the same positive zone of the classifier
- The second one emphasizes the role of the bias as a threshold value to “activate” the +1 output of the classifier.

# Use it: Example (AIMA)


## Classify a new data $(x_1, x_2)$

Find  $h$  s.t. given  $(x_1, x_2)$  return 0/-1 for *Earthquakes* and 1 for *Nuclear Explosion*

Some alg. finds this decision boundary:

$$w_0 + w_1x_1 + w_2x_2 = 0$$

$$-4.9 + 1.7x_1 - 1x_2 = 0$$



hence  $h(6,3) = \text{sign}(w_0 + w_1 \cdot 6 + w_2 \cdot 3) =$   
 $= \text{sign}(-4.9 + 1.7 \cdot 6 - 1 \cdot 3)$   
 $= \text{sign}(2.3) = +1 \rightarrow \text{nuclear expl.}$



*Seismic data.*

$x_1$  body wave magnitude,  $x_2$  surface wave magnitude

Earthquakes (white) Nuclear Explosion (black) 1982-1990 Asia

Getting new examples is expensive;-)



# Use it: Example (Spam)

- Find  $h(mail) + 1$  for *spam*, -1 *not-spam*
  - Features  $\Phi(mail) =$  words [0/1] or phrases ("free money") [0/1] or length [integer]
    - e.g.  $\phi_k(\mathbf{x}) = \text{contain}(\text{word}_k)$  [*bag of words* representation]
- $\mathbf{w} \rightarrow$  weight contribution of the input features to prediction
  - e.g. positive weight for "free money", negative for ".edu" or "unipi"
- $\mathbf{x}^T \mathbf{w}$  is the weight combination
- $h_w(x)$  provides the threshold to decide spam/not spam

$$h_w(\mathbf{x}) = \text{sign}\left(\sum_k w_k \phi_k(\mathbf{x})\right) > 0 \rightarrow +1 = \text{Spam} !$$




# Some useful properties

(we will use them in future lectures)

$$\mathbf{w}^T \mathbf{x} = w_1 x_1 + w_2 x_2 + w_0 = 0$$

Exercise: Draw  $x_2 = -x_1 w_1/w_2 - w_0/w_2$  with different  $w$  values



A linearly separable problem

- If  $w_0=0$  the line goes through the *origin* of the coordinate system.
- If  $n > 2 \rightarrow$  *hyperplane* (decision boundary)
- *Scaling freedom*: the same decision boundary multiplying  $w$  by  $K$
- $w$  is a vector *orthogonal* to the hyperplane:

Given  $x_a, x_b$  (belonging to the sep. hyperplane):  
 $\mathbf{w}^T \mathbf{x}_a + w_0 = 0$ ;  $\mathbf{w}^T \mathbf{x}_b + w_0 = 0$  (take the diff.)  $\rightarrow$   
 $\mathbf{w}^T (\mathbf{x}_a - \mathbf{x}_b) = 0 \rightarrow$  *orthogonal vectors* (dot prod. 0)

- If it exist, there are many possible hyperplanes separating these points:  
also many ML alg.!!!

# Learning Algorithms

- We are going to introduce 2 learning algorithms for the regression and for the classification task using a linear model, both based on LMS:
  1. A direct approach based on **normal equation** solution
  2. An iterative approach based on **gradient descent**
- We start **redefining the learning problem** and the loss for them (for  $l$  data and multidimensional inputs)

# The learning problem (classification tasks)



Dip. Informatica  
University of Pisa

- **Given** a set of  $l$  training examples  $(\mathbf{x}_p, y_p)$  and a loss function (measure)  $L$

$$y_p = \{0, 1\} \text{ or } y_p = \{-1, +1\}$$

- **Find:** The weight vector  $\mathbf{w}$  that minimizes the expected loss on the training data

$$R_{emp} = \frac{1}{l} \sum_{p=1}^l L(h(\mathbf{x}_p), y_p)$$

- For classification: Using a piecewise constant (over  $\text{sign}(\mathbf{w}^T \mathbf{x})$ ) for the loss can make this a difficult problem.
- Assume we still use the *least squares* (as for the regression case)


$$E(\mathbf{w}) \propto \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w})^2 = \sum_{p=1}^l (y_p - \mathbf{w}^T \mathbf{x}_p)^2$$

# Approximating expected loss by a smooth function (#)

- Initially, we can make the optimization problem easier by replacing the original objective function  $L$  (*0/1 loss*) by a **smooth, differentiable function**. For example, consider the popular *mean squared error (MSE loss)*:

Both losses satisfy the minimization of error (\*),

Let us start with LMS avoiding to introduce combinatorial problems



(\*) Example:  $y$  (target)=1,  
 $h(x)=1$  if  $w^T x > 0$ ,  
 $\rightarrow$  No err. when  $w^T x > 0$   
for the classifier

Hence no classif. error  
minimizing either

- 0/1 loss
- or MSE loss

(solution is for both  
on the right part in the plot)

# Learning (a classifier) by Least Squares



Dip. Informatica  
University of Pisa

- Find optimal values for  $\mathbf{w}$  (for fitting of training (TR) data) by *least squares*:
- Given** a set of  $l$  training examples  $(\mathbf{x}_p, y_p)$ , **find**  $\mathbf{w}$  to *minimize* the residual sum of squares:

$$E(\mathbf{w}) = \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w})^2 = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2$$

Where  $\mathbf{x}_p$  is  $p$ -th input vector,  $y_p$  the output for  $p$ ,  $\mathbf{w}$  free par.,  $l$  num. of examples,  $n$  input dim.

Min error: if  $y_p = 1$  then  $\mathbf{x}_p^T \mathbf{w}$  go toward 1  $\rightarrow$  no class. error ;  
if  $y_p = -1$  then  $\mathbf{x}_p^T \mathbf{w}$  go toward -1  $\rightarrow$  no class. Error

**Note:** in  $E(\mathbf{w})$  we do **not** use  $h(\mathbf{x})$ , as for regression, to hold a continuous differentiable loss (because  $h(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x})$  for classification) !!!

- This is a quadratic function  $\rightarrow$  minimum always exists (but may be not unique) [see course of CM or O4DS (@DSBI)]
- $\mathbf{X}$  is a matrix  $l \times n$  with a row for each input vector  $\mathbf{x}_p$
- Note: The same approach is used for a *regression problem*

# Learning Algorithms

- We will introduce 2 learning algorithms for the regression and for the classification tasks using a linear model, both based on LMS

1. A direct approach based on **normal equation** solution

2. An iterative approach based on **gradient descent**



# Normal equation & direct approach solution

- Differentiate  $E(\mathbf{w})$  with respect to  $\mathbf{w}$ :  
Blackboard or Exercise (a next slide).

Result synthesis:

- In the derivation we find that 
$$\frac{\partial E(\mathbf{w})}{\partial w_j} = -2 \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w}) x_{p,j}$$


- We can get the **normal equation**  
(point with gradient of  $E$  w.r.t  $\mathbf{w} = 0$ ):

$$(\mathbf{X}^T \mathbf{X}) \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

- If  $\mathbf{X}^T \mathbf{X}$  is not singular the unique solution is given by

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} = \mathbf{X}^+ \mathbf{y} \quad \text{`+' Moore-Penrose pseudoinverse (also if } \mathbf{X} \text{ is not invertible)}$$

- Else the solution are infinite (satisfying the normal equation):  
we can choose the *min norm* ( $\mathbf{w}$ ) solution



See also  
**CM** course



# Direct approach by SVD

- The *Singular Value Decomposition (SVD)* can be used for computing the pseudoinverse of a matrix ( $X^+$ )

$$X = U \underset{\substack{\downarrow \\ \text{diagonal}}}{\Sigma} V^T \Rightarrow X^+ = V \Sigma^+ U^T$$

by replacing every nonzero entry by its reciprocal

- Moreover we can apply directly SVD to compute  $\mathbf{w} = X^+ \mathbf{y}$  obtaining the minimal norm (on  $\mathbf{w}$ ) solution of least squares problem.
- Note:** THIS IS the learning alg. for the direct approach solution on  $\mathbf{w}$  (or say in *closed form*)
- A practical tool. E.g. see “numerical recipes” in C and many numerical/statistical tools and scientific library (also in R, Octave, Matlab,...)
  - e.g. ARMADILLO: since 2011/12 C++ linear algebra library, NumPy etc.
- Many algorithms addressing the problems of efficiency and stability
- See also CM course (may be later) or O4DS (@DSBI)

# Solution (to find the normal eq.)

## Make as an Exercise !!!(#)



Dip. Informatica  
University of Pisa

$$\begin{aligned}
 \frac{\partial E(\mathbf{w})}{\partial w_j} &= \frac{\partial \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w})^2}{\partial w_j} = \sum_{p=1}^l 2(y_p - \mathbf{x}_p^T \mathbf{w}) \frac{\partial (y_p - \mathbf{x}_p^T \mathbf{w})}{\partial w_j} = \\
 &= \sum_{p=1}^l 2(y_p - \mathbf{x}_p^T \mathbf{w}) \left( 0 - \cancel{\frac{\partial (x_p)_1 w_1}{\partial w_j}} - \cancel{\frac{\partial (x_p)_2 w_2}{\partial w_j}} - \dots \frac{\partial (x_p)_j w_j}{\partial w_j} - \dots \right) = \\
 &= \sum_{p=1}^l 2(y_p - \mathbf{x}_p^T \mathbf{w}) \left( -\frac{\partial (x_p)_j w_j}{\partial w_j} \right) = -2 \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w}) (x_p)_j
 \end{aligned}$$

Only the  
component  $j$   
is not 0

Imposing this =0, we can easily obtain the *normal equation* (first by “sums”, then in matrix notations)

And we also obtained the gradient of  $E$

rewritten as:  $\frac{\partial E(\mathbf{w})}{\partial w_j} = -2 \sum_{p=1}^l \left( (y_p - \mathbf{x}_p^T \mathbf{w}) \right) x_{p,j} = -2 \sum_{p=1}^l \delta_p x_{p,j}$

(Note: In the original image,  $\delta_p$  is written in blue and the term  $(y_p - \mathbf{x}_p^T \mathbf{w})$  is circled with a dashed line.)

(we will use this form in the future, with Neural Networks)

# Other approaches to LS

Many, for instances by an **iterative/gradient descent technique** we can search for:

- More efficient solutions (previous is cubic with dim of matrix  $X$ )
- Regularization (to reduce complexity of the model)
- Better approximation with noisy data by stopping searching before the minimum
- Other approaches that can be applied also to NON-LINEAR models !!!

Hence, we are going to present the second (more important) learning algorithm

# Learning Algorithms

- We will introduce 2 learning algorithms for the regression and for the classification tasks using a linear model, both based on LMS
  1. A direct approach based on **normal equation** solution
  2. An iterative approach based on **gradient descent**

This approach will be the basis for fundamental approaches we will see later




# REPETITA

## Gradient on a surface

- The gradient at a point is a vector pointing in the **direction of the steepest slope** at that point.
- The steepness of the slope at that point is given by the magnitude of the gradient vector
  - The **gradient** shows the direction where the function grows

– The **negative of the gradient** shows the direction where the function *decreases*



# Gradient descent

- Previous derivation suggest the line to construct an iterative algorithm based on:


$$\frac{\partial E(\mathbf{w})}{\partial w_j} = -2 \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w})(\mathbf{x}_p)_j$$

Where  $\mathbf{x}_p$  is  $p$ -th input pattern,  
 $y_p$  the output for  $p$ ,  $\mathbf{w}$  free par.,  
 $l$  num. of examples

→ Component  $j$  of pattern  $p$ , also  $x_{p,j}$

- Gradient = ascent direction:** we can move toward the minimum with a gradient **descent** (changing  $\mathbf{w}$  with  $\Delta \mathbf{w} = -$  gradient of  $E(\mathbf{w})$ )


- *Local search:* it begins with an initial weight vector. We modify it iteratively to decrease up to minimize the error function (steepest descent).



A single  $w$  at step 0,1,2,3,4

# Error surface for linear model with 2 weights ( $\mathbf{w}$ )

Parabolic for the  
 $E(\mathbf{w}) = E([w_0, w_1]^T)$   
(convex quadratic function)



Our "compass" to find the minimum

Hypothesis space  
with 2 parameters

$w_0, w_1$

# The gradient vector

$$\Delta \mathbf{w} = - \frac{\partial E(\mathbf{w})}{\partial \mathbf{w}} = \begin{bmatrix} -\frac{\partial E(\mathbf{w})}{\partial w_0} \\ -\frac{\partial E(\mathbf{w})}{\partial w_1} \\ -\frac{\partial E(\mathbf{w})}{\partial w_2} \\ -\frac{\partial E(\mathbf{w})}{\partial w_j} \\ \dots \\ -\frac{\partial E(\mathbf{w})}{\partial w_n} \end{bmatrix} = \begin{bmatrix} \Delta w_0 \\ \Delta w_1 \\ \Delta w_2 \\ \Delta w_j \\ \dots \\ \Delta w_n \end{bmatrix}$$

We can work in a multi-dim space without the need to visualize it



# Using the Delta Rule

- Hence, as iterative approach we will move using a learning rule based on a «delta» (changing) of  $w$  proportional to the (opposite) of the local gradient
- The «movements» will be made iteratively according to

$$w_{\text{new}} = w + \text{eta} * \Delta w$$

(or component-wise, i.e. for each  $w_j$ )

- that is the “**learning rule**”
- and  $\text{eta}$  ( $\eta$ ) is the “step size” (learning rate) parameter (ruling the speed of our gradient descending)

# Gradient descent algorithm

A simple algorithm:

1) Start with weight vector  $\mathbf{w}_{\text{initial}}$  (small), fix  $\eta$  ( $0 < \eta < 1$ ).

2) Compute  $\Delta \mathbf{w} = -\text{"gradient of } E(\mathbf{w}) \text{"} = - \frac{\partial E(\mathbf{w})}{\partial \mathbf{w}}$  (or for each  $w_j$ )

3) Compute  $\mathbf{w}_{\text{new}} = \mathbf{w} + \eta * \Delta \mathbf{w}$  (or for each  $w_j$ )

Repeat (2) until convergence or  $E(\mathbf{w})$  is "sufficiently small"

- $\Delta \mathbf{w} / l$  : *least mean squares* (dividing by  $l$ , that will be the standard case)

- Batch versions ( $\Delta \mathbf{w}$  after each "epoch" of  $l$  training patterns)

- $\eta$  ( $\eta$ ): step size = *learning rate*: speed/stability trade-off: can be (gradually) decreased to zero (guarantee convergence, avoiding oscillation around the min.): many variants will be introduced later



# Batch/On-line

- For **batch version** the gradient is the sum over all the  $l$  patterns:

$$\frac{\partial E(\mathbf{w})}{\partial w_j} = -2 \left( \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w}) x_{p,j} \right)$$

- provide a more “precise” evaluation of the gradient over a set of  $l$  data

And we upgrade the weights after this sum

- For the **on-line/stochastic version** we upgrade the weights with the error that is computed for each pattern
  - hence, the 2<sup>nd</sup> pattern output is based on weights already updated from the 1<sup>st</sup>, and so ahead
  - It makes progress with each examples it looks at: it can be the faster, but need smaller *eta*:

$$\frac{\partial E_p(\mathbf{w})}{\partial w_j} = -2(y_p - \mathbf{x}_p^T \mathbf{w}) x_{p,j} = -\Delta_p w_j$$

- We will see intermediate cases later (as **mini-batch**)



# Examples

## *Batch algorithm*

We update  $\mathbf{w}$  after (repeating) an "epoch" of  $l$  training data  $\rightarrow$  (blue)


## *On-line algorithm* (stochastic gradient descent - *SGD*)

We update  $\mathbf{w}$  after each pattern  $p$  ( $\Delta_p \mathbf{w}$  for each pattern  $\rightarrow$  (purple and green)



Paths over the error surface by Batch or On-line version

# Learning curve examples



These are **learning curves**:  
They show how the error decreases through gradient descent iterations

P.S. No relation with color in the previous slide!

**Exercise:** *1 is slow, 1 is unstable, 1 is good: which one?*  
*Which with high or low eta value?*

# Gradient descent as Error correction delta rule


$$\Delta w_j = 2 \sum_{p=1}^l (y_p - x_p^T \mathbf{w}) x_{p,j}$$

$\mathbf{w}_{\text{new}} = \mathbf{w} + \text{eta} * \Delta \mathbf{w}$

Where  $x_{p,j}$  is the component  $j$  of the input pattern  $p$ ,  $y_p$  the output for  $p$ ,  $\mathbf{w}$  free par.,  $l$  num. of examples.


*The constant 2 can be omitted*

- This is an "error correction" rule (Widrow-Hoff or **delta rule**) that change each  $w_j$  proportionally to the error (target  $y$  - output):
  - E.g. (target  $y$  – output) = err=0  $\rightarrow$  no correction
  - (input $_j$ >0) if err + (output is too low), positive delta  $\rightarrow$  increase  $w_j \rightarrow$  increment the output  $\rightarrow$  less err
  - (input $_j$ >0) if err - (output is too high), negative delta  $\rightarrow$  decrease  $w_j \rightarrow$  reduce output  $\rightarrow$  less err
  - ... [exercise: all the cases]
- We improve by learning from previous errors  
"seeking and blundering we learn (Goethe)"




# Delta-W as Error Correction Learning rule (II)

Inputs      Weights      ( $w_0=0$ )       $\Delta w_j = 2 \sum_{p=1}^l (y_p - x_p^T \mathbf{w}) x_{p,j}$



If **misclassified** (because target is  $+1$ )  $\rightarrow (1 - (???)) \rightarrow$  High positive delta for  $w_1$  and  $w_3 \rightarrow$  increase them proportionally (with eta) to the *delta*, hence a positive value in this case  
[error correction rule] !!!



e.g. (see figure):

Now is **correct** !!!

Exercise: compute the values for delta (hint: see eq. above) and for also eta

# Gradient descent: final discussion



Dip. Informatica  
University of Pisa

**Gradient descent** approach it is simple and effective *local search* approach to LMS solution and

- It allows us to search through an *infinite (continuous) hypothesis* space!
- It can be easily always applied for *continues H* and *differentiable loss*
- NOT ONLY to linear models !!!! (we will see for Neural Networks and deep learning models)
- *Efficient?* Many improvements are possible, e.g. Newton & quasi-newton methods; Conjugate Gradient, ...! → [CM course\\*](#) (we will mention examples later, discussing Neural Networks)



# Summarizing (Classification)

- Model **trained** (on TR set) with LS (LMS) on  $\mathbf{w}^T \mathbf{x}$ 
  - by the simple gradient descent algorithm used for linear regression
- Model **used** for classification applying a threshold function, obtaining  $h(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x})$
- The error can be computed as *classification error* or number of misclassified patterns (not only by the Mean Square Error)

0/1 Loss

$$L(h(\mathbf{x}_p), d_p) = \begin{cases} 0 & \text{if } h(\mathbf{x}_p) = d_p \\ 1 & \text{otherwise} \end{cases}$$

$$\text{mean\_err} = \frac{1}{l} \sum_{p=1}^l L(h(\mathbf{x}_p), d_p)$$

*num\_err*

- ACCURACY** = mean of correctly classified =  $(l - \text{num\_err}) / l$

# Coming back to the problem: Linear model solution

Linear Regression of 0/1 Response




Figure 2.1 (© HTF 2001):

*A classification example in two dimensions. The classes are coded as a binary variable GREEN = 0, RED = 1 and then fit by linear regression. The line is the decision boundary defined by  $\mathbf{x}^T \mathbf{w} = 0.5$ . The red shaded region denotes that part of input space classified as RED, while the green region is classified as GREEN.*

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{x}^T \mathbf{w} > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

Linear threshold unit

The *decision boundary* is  $\{\mathbf{x} \mid \mathbf{x}^T \mathbf{w} = 0.5\}$  is linear (and seems to make many errors on the training data). Is it true?

# Good or bad approximation?

(#)



Dip. Informatica  
University of Pisa

- Possible scenarios (we know the true target function!)
  - Scenario 1: The data in each class are generated from a Gaussian distribution with uncorrelated components, same variances, and different means.
  - Scenario 2: The data in each class are generated from a mixture of 10 gaussians in each class.
- For Scenario 1, the linear regression rule (by LS) is almost **optimal** (is the best one can do). The region of overlap is inevitable (due to errors in the input data).
- For Scenario 2, it is far too rigid: next models for it!

\* Least squares corresponds to the maximum likelihood criterion if the experimental errors have a normal distribution

# Linear model (in ML): Inductive Bias (alla Mitchell)




Dip. Informatica  
University of Pisa

- **Language bias**: the  $H$  is a set of linear functions (may be very restrictive and rigid)
- **Search bias**: ordered search guided by the Least Squares minimization goal
  - For instance, we could prefer a different method to obtain a restriction on the values of parameters, achieving a different solutions with other properties (in particular to consider the generalization issue), ...

It shows that even for a “simple” model there are many possibilities.  
We need a principled approach! (see theory of ML)...

# Limitations: regression tasks for non linear problems




Too poor solution

# Limitations (classification) (language bias)



Dip. Informatica  
University of Pisa

- In geometry, two set of points in a two-dimensional plot are **linearly separable** when the two sets of points can be completely separated by a single line
- In general, two groups are *linearly separable* in  $n$ -dimensional space if they can be separated by an  $(n - 1)$ -dimensional hyperplane.



- The linear decision boundary can provide exact solutions only for linearly separable sets of points

# Example: Conjunctions

- We can represent conjunctions by the linear models, e.g.:
- Conjunctions** (see the example in the introduction lectures):

4 var.:  $x_1 \wedge x_2 \wedge x_4 \leftrightarrow y$



- $1 \ x_1 + 1 \ x_2 + 0 \ x_3 + 1 \ x_4 \geq 2.5$

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w}\mathbf{x} + w_0 \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

In the plot:

2 var.:  $x_1 \wedge x_2 \leftrightarrow y$

- $1 \ x_1 + 1 \ x_2 \geq 1.5$



$\mathbf{w}$  can be learned to find this solution


# Limitations: classification tasks



Dip. Informatica  
University of Pisa









- Given **3 points**, can we always find a separation plane for every assignment of  $f(\mathbf{x})$ ?  
**No, 3 aligned points with 0 in the middle and others 1; yes if they are not aligned (existence!).**

Here NOT complete:  
There are  $2^3$  cases



- Given **4 points**, can we always find a separation plane for every assignment of  $f(\mathbf{x})$ ?  
**No (XOR)**

we can find a labeling such that the linear classifier fails to be perfect


		XOR
		00 0
		01 1
		10 1
		11 0

**Exercise:** try to redo it by yourself!



# How to extend the linear model

- Note that in  $h_{\mathbf{w}}(x) = w_1 x + w_0$  or  $h_{\mathbf{w}}(\mathbf{x}) = \mathbf{w}^T \cdot \mathbf{x}$
- As Statistical Parametric models:  
"linear" does not refer to this (red) straight line, but rather to the way in which the regression coefficients  $\mathbf{w}$  occur in the regression equation



- Hence, we can use also transformed inputs, such as  $x, x^2, x^3, x^4, \dots$  with *non-linear* relationship inputs and output, holding the learning machinery (Least Square solution) developed so far...

$$h_{\mathbf{w}}(x) = w_0 + w_1 x + w_2 x^2 + \dots + w_M x^M = \sum_{j=0}^M w_j x^j$$

polynomial regression

# A generalization (**LBE**) (shown for regression\*)



Dip. Informatica  
University of Pisa

- Basis transformation: **linear basis expansion (LBE)**:

$$h_w(x) = \sum_{k=0}^K w_k \phi_k(x)$$

- Augment the input vector with additional variables which are transformations of  $x$  according to a function  $\phi$  ( $\phi_k: R^n \rightarrow R$ )
- E.g.
  - Polynomial representation of  $x$ :  $\phi(x) = x_j^2$  or  $\phi(x) = x_j x_i$ , or ....
  - Non-linear transformation of single inputs:  $\phi(x) = \log(x_j)$ ,  $\phi(x) = \text{root}(x_j)$ , ....
  - Non-linear transformation of multiple input:  $\phi(x) = ||x||$
  - *Splines*, ..., ....
- Typically: Number parameters  $K > n$  (before it was  $n$ )
- The model is *linear in the parameters (also in  $\phi$ , not in  $x$ ): we can use the **same learning alg.** as before!*
- *Note: it can be applied for regression (here) or classification (Exercise: HOW?)*

# Basis Expansion: examples



Dip. Informatica  
University of Pisa

- Basis transformation: **linear basis expansion**:

$$h_{\mathbf{w}}(\mathbf{x}) = \sum_{k=0}^K w_k \phi_k(\mathbf{x})$$

*EXAMPLES:*

- [1-dim  $\mathbf{x}$ ]  $\phi_j(x) = x^j$ .

$$h(\mathbf{x}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

1-dim polynomial regression ( $K=M$ )

Already seen in  
the introduction

- Or “any other”, e.g.  $\phi(\mathbf{x}) = \phi([x_1, x_2, x_3])$

$$h(\mathbf{x}) = w_1x_1 + w_2x_2 + w_3\log(x_2) + w_4\log(x_3) + w_5(x_2x_3) + w_0$$

# Basis Expansion criticism

- Which  $\phi$  ( $\phi$ ) ? Toward the so called “**dictionary**” approaches
  - **PROS**: Can model more complicated relationships (than linear) w.r.t. the inputs: it is more expressive.
  - **CONS**: With *large* basis of functions, we easily risk *overfitting*, hence we require methods for controlling the complexity (\*) of the model
    - *Curse of dimensionality* (the volume of the problem space increases so fast that the available data become sparse, the data became no more sufficient to support the model complexity)  $\rightarrow$  we will see later
    - $\Phi$  are *fixed before* observing training data
      - versus adaptive /non-linear in parameters e.g. NN!
  - *We will see the alternative NN and SVM solutions!*
- (\*) Whereas **complexity** is **not** for the computational cost but a measure of the flexibility of the model to fit the data (see the VC-dim)

# Improvements: How to control model complexity? **Important !**



Dip. Informatica  
University of Pisa

- Many approaches... e.g. coefficient shrinkage:
- **Ridge regression**: (**Tikhonov regularization**): smoothed model → possible to add constraints to the sum of value of  $|w_j|$  penalizing models with high values of  $|w|$ , i.e. favoring "sparse" models using less terms due to weights  $w_j = 0$  (or close to 0) (it means a less complex model)

"Error" data term [(M)SE]      Regularization/penalty term  
( $R_{emp}$  in the SLT)

$$Loss(\mathbf{w}) = \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w})^2 + \lambda \|\mathbf{w}\|^2 \longrightarrow \sum w_j^2$$

Lambda ( $\lambda$ ): regularization (hyper)parameter  
(a small positive value chosen by the "model selection" phase)

The sum is over  
the number of  $w$

- Note that for the **objective function** we use here the name *Loss* (used for the model training cost function) to distinguish from the *Error*  $E$  (useful to evaluate the model error and used for the *data term* inside this Loss).  
Hence, differently from previous assumptions the two terms are not more equivalent in our use (in the course).

# Tikhonov Regularization: Solving it



Dip. Informatica  
University of Pisa

$$Loss(\mathbf{w}) = \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w})^2 + \lambda ||\mathbf{w}||^2$$

- For the direct approach:

$$\mathbf{w} = \boxed{(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1}} \mathbf{X}^T \mathbf{y} \quad \text{-----} \rightarrow \text{this matrix is always invertible}$$

- For the gradient approach: we still apply –gradient of the Loss:
  - As an **Exercise** compute the gradient of the two terms (error and penalty terms) of the  $Loss(\mathbf{w})$  w.r.t. weights  $w_i$  separately, using *eta* only for term  $E$
  - Formulate the new update rule :
  - You will obtain:  $\mathbf{w}_{\text{new}} = \mathbf{w} + \text{eta} * \Delta \mathbf{w} - 2 \lambda \mathbf{w}$
  - That is a **weight decay** technique (basically add  $2\lambda \mathbf{w}$  to the gradient)
    - E.g. with 0 gradient, it decreases the value of each  $w$  with a fraction of the old  $w$

# Andrej Nikolaevič Tikhonov

---



Russian mathematician  
1906 –1993

# Tikhonov regularization: trade-off



Dip. Informatica  
University of Pisa

$$Loss(\mathbf{w}) = \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w})^2 + \lambda ||\mathbf{w}||^2$$

- Note the balancing (trade-off) between the two terms:
  - Small **lambda** ( $\lambda$ ) value  $\rightarrow$  minimizing the loss the focus is on obtaining a small error data term (first term, minimize just the training error) with a too complex model (high norm of the weights), the risk is of **overfitting**,
  - High **lambda** ( $\lambda$ )  $\rightarrow$  minimizing the loss the focus is on the second term, hence the data error (first term) could grow too much, i.e. moving to **underfitting**
  - The trade-off is ruled by the value of **lambda** ( $\lambda$ )
  - We will see soon some examples
- The main advantage is that we have a concrete realization of the control of model complexity, easy to be implemented and of general applicability.



# Tikhonov and SLT


- The penalty term **penalizes high value of the weights** and tends to drive all the weights to smaller values
  - E.g. Some weights values can go even to zero
- It **implements** a control of the model complexity
- This leads to a model with **less (or proper) *VC-Dim***, with a trade-off obtained through **just a (1) parameter** that you can control: the  $\lambda$
- **Exercise**: connect Tikhonov regularization with the slide on SLT (see the introduction) to see
  - why this can help to have a better bound on  $R$ , and
  - how lambda values can rule the underfitting/overfitting cases
- **Exercise**: derive the new learning rule with weight decay using the Tikhonov loss: computing again the partial derivatives of  $Loss$  with respect to  $w$ , separating data term ( $\bullet$  eta) and penalty term ( $\bullet$   $\lambda$ )

# Slide from Intro slides: 9<sup>th</sup> Order Polynomial



Dip. Informatica  
University of Pisa

As to have Lambda 0



We have already seen it  
 $E(\mathbf{w}) = 0$  on training data and **overfitting**

# Regularization effect

$$\ln \lambda = -18$$




Dip. Informatica  
University of Pisa

- 9<sup>th</sup> Order Polynomial

Lambda small positive


$$\lambda = 0.0000000152$$



Once regularized with a suitable value of lambda it works well!  
It is worse on training data but we reduce the **overfitting**


# Too High lambda ( $\ln \lambda = 0$ )

Very high lambda, close to 1



But we need a trade-off ...  
(because a too high lambda leads as to **underfitting**)

# Regularization: $E_{\text{RMS}}$ vs $\ln \lambda$



Low lambda  $\rightarrow$  overfitting

High lambda  $\rightarrow$  underfitting

# Polynomial Coefficients

	0 lambda ↗		↖ Very high lambda
	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
$w_0^*$	0.35	0.35	0.13
$w_1^*$	232.37	4.74	-0.05
$w_2^*$	-5321.83	-0.77	-0.06
$w_3^*$	48568.31	-31.97	-0.05
$w_4^*$	-231639.30	-3.89	-0.03
$w_5^*$	640042.26	55.28	-0.02
$w_6^*$	-1061800.52	41.32	-0.01
$w_7^*$	1042400.18	-45.95	-0.00
$w_8^*$	-557682.99	-91.53	0.00
$w_9^*$	125201.43	72.68	0.01

# Other Regularization Tech. for Linear Models



Dip. Informatica  
University of Pisa

- Ridge regression ( $\| \cdot \|_2$ )
  - Lasso ( $\| \cdot \|_1$ )
  - Elastic nets (use both  $\| \cdot \|_1$  and  $\| \cdot \|_2$ )
- 
- We will introduce them later in the course (or in CM course), anyway:
  - The *L2 norm* penalizes the square value of the weight and tends to drive all the weights to smaller values.
  - On the other hand, the *L1 norm* penalizes the absolute value of the weights and tends to drive some weights to exactly zero (while allowing some weights to be large) → toward feature selection!
    - Unfortunately  $\| \cdot \|_1$  (absolute value) introduce a non differentiable loss → needs other approaches



# Other Improvements (optional)



Dip. Informatica  
University of Pisa

- Data augmentation:
  - *Inputs with added noise*: teaching models to ignore irrelevant variations, it contributes to robustness and *regularization*
  - Oversampling to address imbalanced datasets
  - ...
- Derived inputs

a small number of new variables is used in place of the  $\mathbf{x}$  inputs, which are a linear combination of  $\mathbf{x}$ .

  - Principal Component Regression
  - Partial Least Squares





# Multi-class task (#) (optional)



Dip. Informatica  
University of Pisa

Two very simple approaches for multi-class:

- Class 1-of-K rep: {red, green, blue}  $\rightarrow$  (0,0,1), (0,1,0), (1,0,0).  $\rightarrow$  solve 3 linear models
- **OVA**: "one-vs-all": a discriminant function for each class

built on top of real-valued binary classifiers


- train K different binary classifiers, each one trained to distinguish the examples in a single class from the examples in all remaining classes.
- to classify a new example, the K classifiers are run, and the classifier which outputs the largest (most positive) value is chosen.
- **AVA**: "all-vs-all" = "one-versus-one":
  - each classifier separates a pair of classes. Apply it to all the pairs:  $K(K-1)$   
[How many training?]
  - to classify a new example, all the classifiers are run, and the winner is the one with the max sum of outputs versus all the other classes **OR** the class with most votes
  - training data set for each classifier is much smaller

# Criticism (##) (optional)



Dip. Informatica  
University of Pisa

The problem can become not easy



- Masking: classes can be masked by others (for high K)
- A wide array of more sophisticated approaches for multiclass classification exists ...
- Some models can deal directly with multi-output

# Other learner models for classification (#) (optional)



Dip. Informatica  
University of Pisa

- Linear Discriminant Analysis (also multi-class)
- Logistic regression
  - $P(y|\mathbf{x})$  starting from modeling the class density as a known density

# Extensions: ML Models

## (pro future) (#)



Dip. Informatica  
University of Pisa

In ML course (read it later!):

- Perceptron (Rosenblatt 1958, biological inspiration):
  - “minimize (only) misclassifications” algorithm
  - basis for Neural Networks (set of units with layers)
  - Adaptive *basis expansions* (*phi learned by training*)
    - Feature representation learning in each layer (deep learning concept)
  - Gradient descent approach for learning
- SVM (Vapnik 1996):
  - *regularization* via the concept of maximum margin: Maximize the gap (margin) between the two classes on the training data.
  - enlarge the feature space via *basis expansions* (e.g. polynomials).
- NN and SVM models realize (also) a flexible ***non-linear*** approximation for classification and regression problems

# Lessons Learned

- We concretely showed that is possible to *learn* by tuning free parameters  $\mathbf{w}$ , searching over a continuous space guided by a loss  $f$ .
  - How to state/formulate a regression/classification learning problem for a linear model by LMS
  - How to derive the learning algorithm (basic univariate, direct and **gradient** based forms)
  - Two learning algorithms
  - How the model works after training
  - Linear models have limitations (strong language bias)
- How to extend to non-linear modeling (by **LBE**)
- How to implement a **regularization approach** changing the loss  $f$ .
- The role of regularization for the **control of the model complexity**

Also useful for your self assessment (try to check if you can answer)

# ML Course structure

## Discussion on Linear Model

---

In the *file rouge* of ML, 2 main concepts up to now:

- Basis expansion → (implement) more flexibility
- Regularization → (implement) control of complexity
- The cases of the *linear models* provided an *instance* routed in the classical mathematical approaches but the we will find again these concepts, by different models and implemented in different/similar forms, in the modern ML (e.g. for NN & SVM in our course)!
- Now we move on the other extreme, toward a very flexible approach (K-nn)

# Bibliography: Linear and K-nn (I)



Dip. Informatica  
University of Pisa

- Hastie, Tibshirani, Friedman, The Elements of Statistical Learning, Springer Verlag, 2001-2017 (*check the new Ed.s*): **chap 2**  
(note: it exists an on-line pdf version)
- Mitchell:
  - Linear model and LMS alg.: chap 4.4
  - K-nn: **chap 8**
- Haykin (2nd edition):
  - Linear model and LMS alg.: chap 3 up to 3.7  
[details on Newton and Gauss-Newton methods are not strictly needed]
- Haykin (3rd edition):
  - Linear model and LMS alg.: chap 3  
[details on Newton and Gauss-Newton and other advanced approaches are not strictly needed]

# Bibliography: Linear and K-nn (II)



Dip. Informatica  
University of Pisa

- Moreover...
- Further readings (not mandatory!):
  - Gently introduction to linear models:  
AIMA (Russel, Norvig, Artificial Intelligence: A Modern Approach), ed .3: **chap 18.6** (18.6.1,18.6.2,18.6.3)

- To go in deep for **Linear least squares**

You can start from www resource as (With nice examples):

[https://en.wikipedia.org/wiki/Linear\\_least\\_squares\\_%28mathematics%29](https://en.wikipedia.org/wiki/Linear_least_squares_%28mathematics%29)

- **LS in general**

[http://en.wikipedia.org/wiki/Least\\_squares](http://en.wikipedia.org/wiki/Least_squares)



# For information

**Alessio Micheli**  
**micheli@di.unipi.it**

[www.di.unipi.it/groups/ciml](http://www.di.unipi.it/groups/ciml)



Dipartimento di Informatica  
Università di Pisa - Italy  
Micheli



**Computational Intelligence &  
Machine Learning Group**