

# K-nn

---

**Alessio Micheli**

**[micheli@di.unipi.it](mailto:micheli@di.unipi.it)**



Dipartimento di Informatica  
Università di Pisa - Italy

**Computational Intelligence &  
Machine Learning Group**

*Sept, 2025*

# K-nn

---

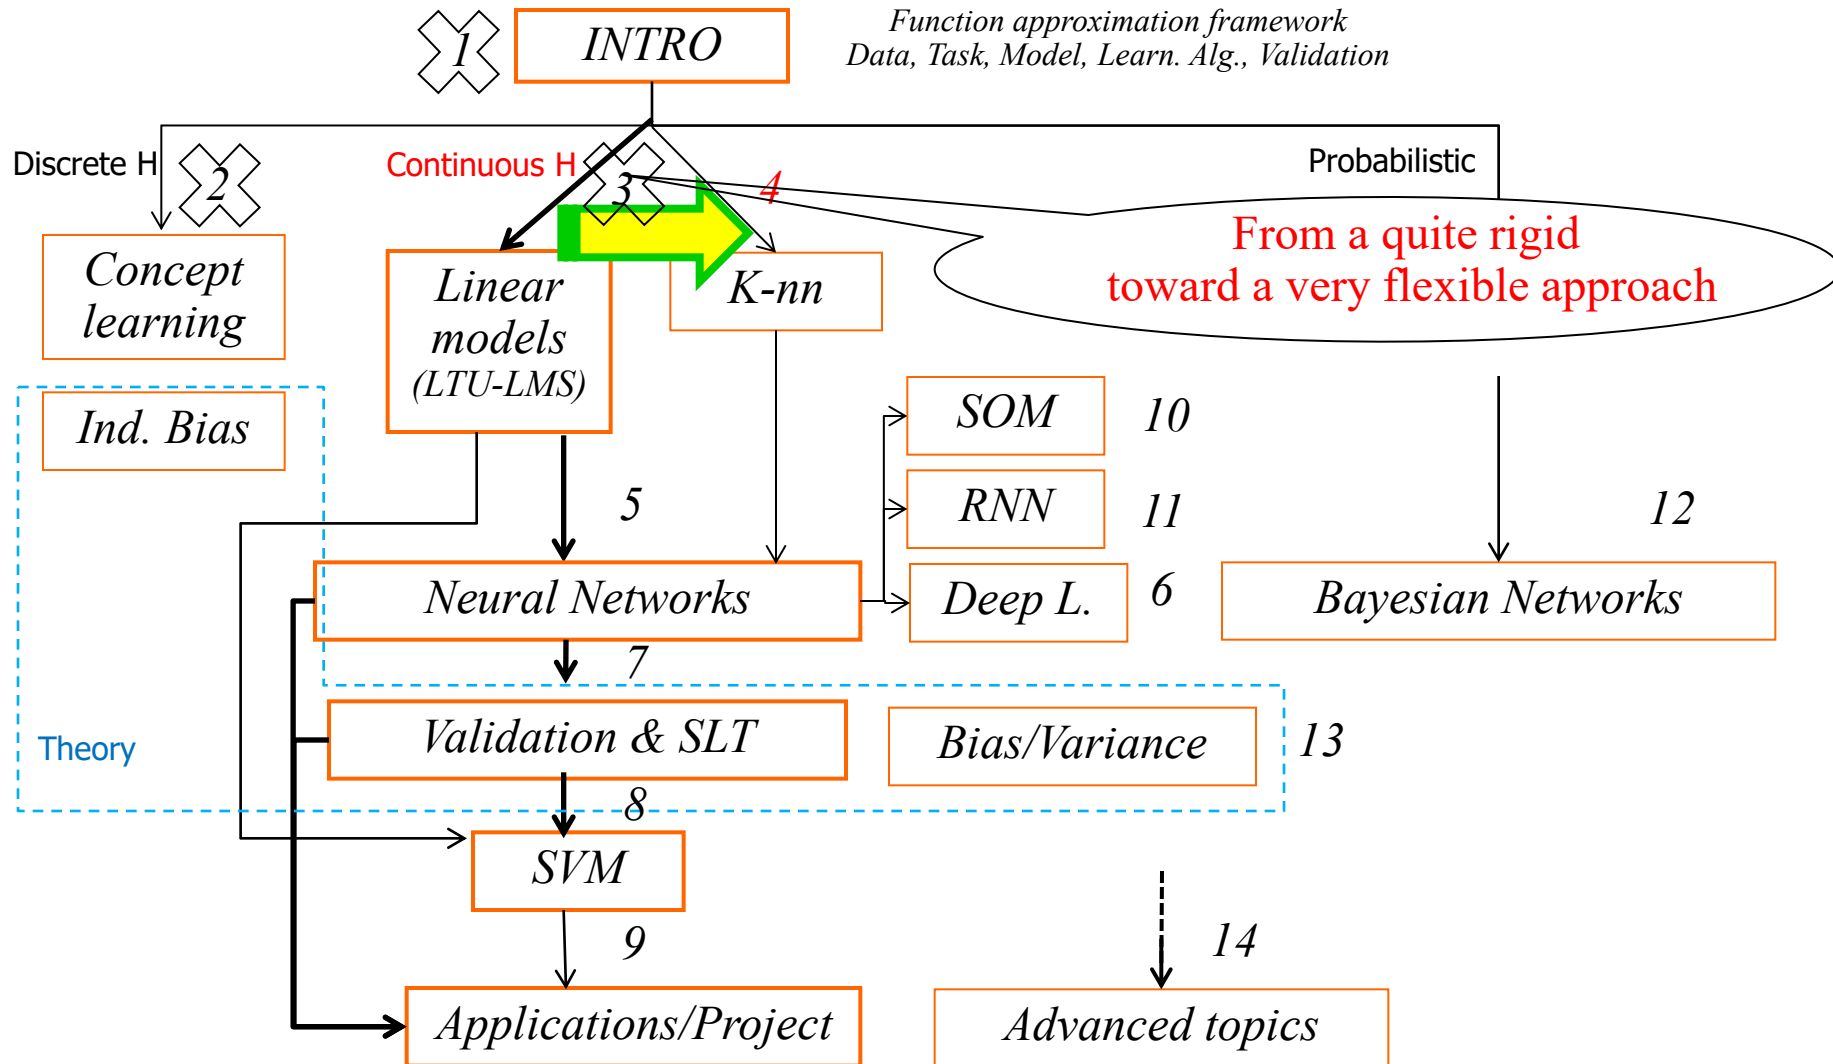
**Still simple but flexible (and local)**

# ML Course structure

## Where we go




Dip. Informatica  
University of Pisa



# Repetita: Problem: example

Raw Data with a Binary Response



200 points generated in  $\mathbb{R}^2$  from an unknown distribution; 100 in each of two classes.

Can we build a rule to predict the color of future points?

Data may be generated by gaussian distribution (for each class) with different means  
or by a mixture of different low variance gaussian distributions.

# Repetita: Find by LS a line to separate

Linear Regression of 0/1 Response




Figure 2.1 (© HTF 2001):

*A classification example in two dimensions. The classes are coded as a binary variable GREEN = 0, RED = 1 and then fit by linear regression. The line is the decision boundary defined by  $\mathbf{x}^T \mathbf{w} = 0.5$ . The red shaded region denotes that part of input space classified as RED, while the green region is classified as GREEN.*

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{x}^T \mathbf{w} > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

Linear threshold unit

The decision boundary is  $\{\mathbf{x} \mid \mathbf{x}^T \mathbf{w} = 0.5\}$  is linear (and seems to make many errors on the training data). Is it true?

# Learning... timing

- **LEARNING ALG.**
- Timing.
  - **Eager:** Analyze the training data and construct an explicit hypothesis.
  - **Lazy:** Store the training data and wait until a test data point is presented, then construct an ad hoc hypothesis to classify that one data point.

# 1-Nearest Neighbor (1-nn)

The algorithm:

- Simply store the training data  $\langle \mathbf{x}_p, y_p \rangle$   $p=1 \dots l$
- Given an input  $\mathbf{x}$  (with dimension  $n$ )
- Find the nearest training example  $\mathbf{x}_i$ 
  - Find  $i$  s.t. we have  $\min d(\mathbf{x}, \mathbf{x}_i) \rightarrow i(\mathbf{x}) = \arg \min_p d(\mathbf{x}, \mathbf{x}_p)$
  - E.g. Euclidian distance :
- Then output  $y_i$

$$d(\mathbf{x}, \mathbf{x}_p) = \sqrt{\sum_{t=1}^n (x_t - x_{p,t})^2} = ||\mathbf{x} - \mathbf{x}_p||$$

↓  
Pattern  $\mathbf{x}_p$ , component  $t$

# 1-nn

1-Nearest Neighbor Classifier




Figure 2.3:

*The same classification example in two dimensions as in Figure 2.1.*

GREEN = 0, RED = 1


- Very flexible !
- No misclassifications in TR data: 0 training error: what for test data ?
- Decision boundaries is *not* linear: it is quite *irregular*
- May be unnecessary noisy (e.g. for scenario 1)



# 1-nn vs 5-nn example

- 1-nn return + for  $x_q$
- 5-nn return – for  $x_q$

**Smoothing** over a set of neighbors for noisy data



# K-Nearest Neighbors

- A natural way to classify a new point is to have a look at its neighbors,
- and take an average:

$$avg_k(\mathbf{x}) = 1/k \sum_{\mathbf{x}_i \in N_k(\mathbf{x})} y_i$$

- where  $N_k(\mathbf{x})$  is a neighborhood of  $\mathbf{x}$  that contains exactly  $k$  neighbors (closest patterns according to *distance*  $d$ ):
  - k-nearest neighborhood: **K-nn**.
- If there is a clear dominance of one of the classes in the neighborhood of an observation  $\mathbf{x}$ , then it is likely that the observation itself would belong to that class, too. Thus the classification rule is the **majority voting** among the members of  $N_k(\mathbf{x})$ . As before,

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } avg_k(\mathbf{x}) > 0.5 \\ 0 & \text{otherwise} \end{cases} \quad \text{for targets } y \text{ in } \{0,1\}$$

- For regression task: use directly the *avg*: mean over K-nn

# 15-nn

15-Nearest Neighbor Classifier




Figure 2.2:


*The same classification example in two dimensions as in Figure 2.1.*

**GREEN** = 0, **RED** = 1

*The predicted class is hence chosen by majority vote amongst the 15-nearest neighbors.*

- Still very flexible !
- Some misclassifications in TR data
- Decision boundaries is *not* linear: it is still quite, although less, *irregular*
- Decision boundary adapts to the local densities of the classes

# Voronoi diagram



Each cell consisting of all points closer to  $x$  than to any other patterns.

The segments of the Voronoi diagram are all the points in the plane that are equidistant to two patterns.

**Implicitly used by K-nn**



# K-nn for multi-class

- Return the class most common amongst its  $k$  nearest neighbors

$$h(\mathbf{x}) = \arg \max_v \sum_{\mathbf{x}_i \in N_k(\mathbf{x})} \mathbf{1}_{v, y_i}$$

$$\mathbf{1}_{v, y_i} = \begin{cases} 1 & \text{if } v = y_i \\ 0 & \text{otherwise} \end{cases}$$

We count the classes in the neighbor (by the  $\mathbf{1}_{v, y}$  for each  $v$ )  
taking the most frequent class ( $\arg \max$ )

# K-nn variants: Weighted distance



Dip. Informatica  
University of Pisa

- It can be useful to weight the contributions of the neighbors, so that the nearer neighbors contribute more to the average than the more distant ones.

$$h(\mathbf{x}) = \arg \max_v \sum_{\mathbf{x}_i \in N_k(\mathbf{x})} \mathbf{1}_{v, y_i} \cdot \frac{1}{d(\mathbf{x}, \mathbf{x}_i)^2}$$

If  $d = 0$  for a  $i$ , return  $y_i$

$$\mathbf{1}_{v, y_i} = \begin{cases} 1 & \text{if } v = y_i \\ 0 & \text{otherwise} \end{cases}$$

# K-nn: An extreme

- Not a global hypothesis for all the instances → no model to be fit
  - We need to memorize the input examples
- Local estimations (by locally constant functions) vs global linear approximation/estimation of the target function (over the instance space)
- Lazy, memory based, instance-based, **distance-based methods**

# K-nn versus Linear

## Discussion on linear versus k-nn models

Two extremes of the ML panorama:

- Rigid (low variance) versus flexible (high variance):
  - In K-nn with small  $k$  few points can change the decision boundary
  - We may pay a price for this flexibility
- Eager versus lazy
- Parametric versus instance-based
- Linear regression uses 3 parameters to describe its fit in the example of Fig. 2.1,  $w_0, w_1, w_2$  (or  $n+1$  in general)  
Does K-nn use 1 (the value of  $k$  here)?
- More realistically, K-nn uses  $l/k$  “effective number of parameters” (Hastie-Tibshirani-Friedman 2001)\*
  - where  $l$  is used the number of data ( *$N$  in the book*)



# LS linear vs K-nn with various k values

Figure 2.4: *Misclassification curves for the simulation example used in Fig. 2.1, 2.2 and 2.3.*

Training set of size 200

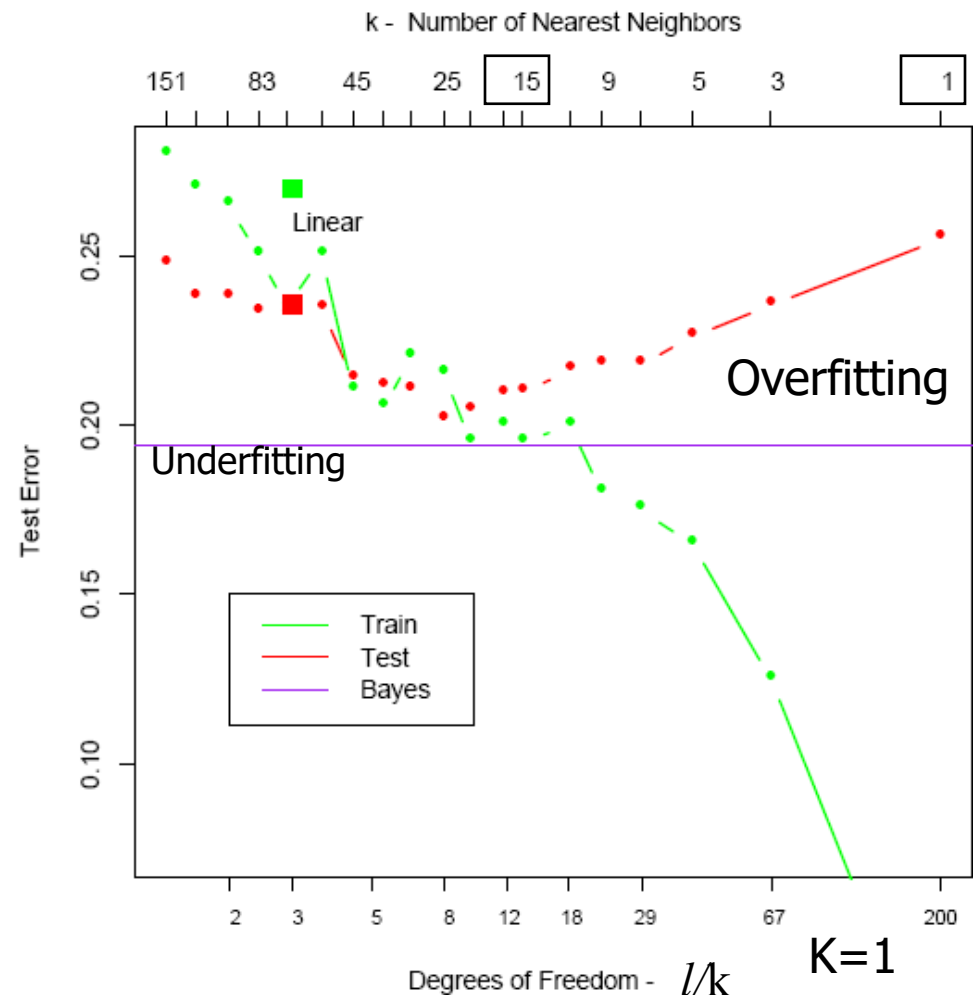
Test set of size 10.000.

The **red** curves are test and the **green** are training error for k-nn classification (changing K).

The results for linear regression are the bigger green and red dots at three degrees of freedom.

The **purple** line is the optimal Bayes Error Rate (See next slides)

Note how we move from underfitting to overfitting moving the values of k (i.e. the rate  $l/k$ ): more flexibility allows to find the best result if we control “complexity” by K

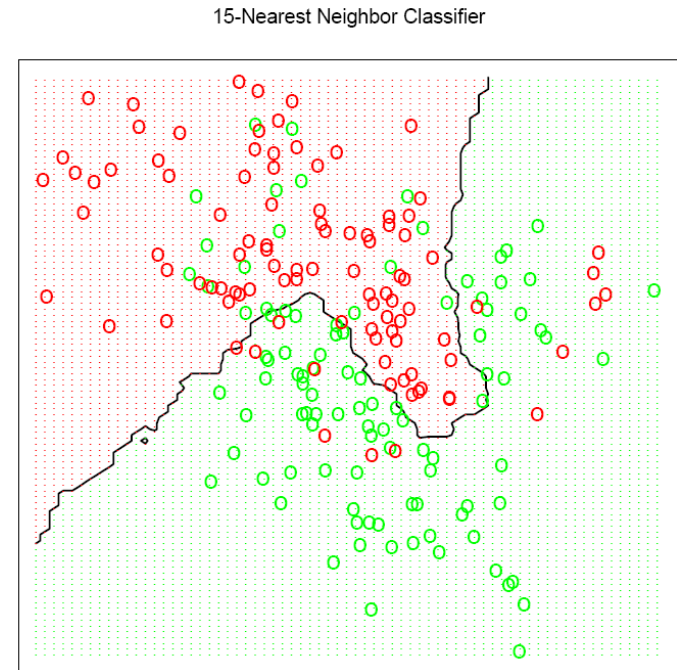



# Bayes error rate

Bayes optimal solution (called Bayes classifier):

- If we know the density  $P(x,y)$  we classify to the most probable class, using the conditional (discrete) distribution as:
  - Output the class  $v$  s.t. is  $\max P(v|x)$  ( $v$  in  $\{C_1, C_2, C_3 \dots C_K\}$ )
- The error rate of the (optimal) Bayes classifier is called the Bayes rate.
- I.e. the minimum achievable error rate given the distribution of the data (assuming that the generating density is known !!!)
- Note on K-nn: we see that the k-nn classifier directly *approximates this solution* (a majority vote in a nearest neighborhood amounts to exactly this), except that
  - conditional probability at a point is relaxed to conditional probability within a neighborhood of a point (local approximation), and
  - probabilities are estimated by training-sample proportions.

# Bayes Optimal Classifier: results



- The optimal Bayes decision boundary for the simulation example.
- Since the generating density is known for each class, this boundary can be calculated exactly
- 15-nn was *close to this* (indeed it shown a min. test err in the plot a couple of slides ago)

# Inductive Bias of k-nn


1. The assumed **distance** tells us which are the most similar examples
  - The classification is assumed similar to the classification of the neighbors according to the assumed metric
    - E.g. we can use other metric than the Euclidian
    - Symbolic data require “ad-hoc” metrics (e.g. Hamming distance between two strings of equal length is the number of positions for which the corresponding symbols are different).
2. **Locality**: local approximation (local smoothness prior): see before, and later in the course (vs Deep Learning bias)
  - Next slides:
    - criticism and limitations

# 1. Scale changes and 2. other metrics



Dip. Informatica  
University of Pisa

- **Scale changes:** Domain knowledge dependent choice
- **If** variables should contribute equally:
  - Pay attention to disparity in the ranges of each variables
  - Rescale data to equalize inputs ranges ( $\rightarrow$  change the metric)!  
E.g. mean zero and variance 1 normalization
- Variable scaling can have a high impact (i.e. k-nn is fragile even with respect to basic preprocessing)



# (some) Limits of K-nn: computational cost



Dip. Informatica  
University of Pisa

- Note that K-nn makes the local approximation to the target function for each new example to be predicted:
  - The computational cost is deferred to the **prediction phase!**

Moreover: high retrieval cost:

- Computationally intensive (in time) for each new input: computing the distances from the test sample *to all* stored vectors
  - The time is proportional to the number of stored patterns
  - “ad-hoc” proximity search algorithms to optimize
  - E.g. by indexing the patterns or other approximate nearest neighbor search techniques
- Cost in space (all the training data)

# **(some) Limits of K-nn**

---



Dip. Informatica  
University of Pisa

- K-nn models offer little interpretation.
  - Subjectivity of interpretation
  - Dependence on the metric

# (some) Limits of K-nn: curse of dim.



Dip. Informatica  
University of Pisa

K-nn provides a good approximation if we can find a significant set of data close to any  $\mathbf{x}$ , with dense sampling

## When can it fail?

- When we have a lot of input variables (high  $n$ , *i.e.* high dim.), K-nn methods often fails because of the **curse of dimensionality**:

Many manifestations of this problem: we will examine a few

1. *It is hard to find nearby ("similar") points in high dimensions!*
2. *Low sampling density for high-dim data*
3. *Irrelevant features issue*

*These are important in terms of ML,  
as they inherently affect the **generalization capability** !!!*



# (some) Limits of K-nn: curse of dim.



Dip. Informatica  
University of Pisa

## *1. It is hard to find nearby points in high dimensions!*

- K-nearest neighbors can fail in high dimensions, because it becomes difficult to gather K observations “close” (in terms of variables values, i.e. “similar”) to a query point  $x_q$ :
  - near neighborhoods tend to be **spatially large**, and estimates are **not longer local**
  - reducing the spatial size of the neighborhood means reducing K, and the variance of the estimate increases ( $\rightarrow$  overfitting).
- Why? See the next slide...

# (some) Limits of K-nn: curse of dim.

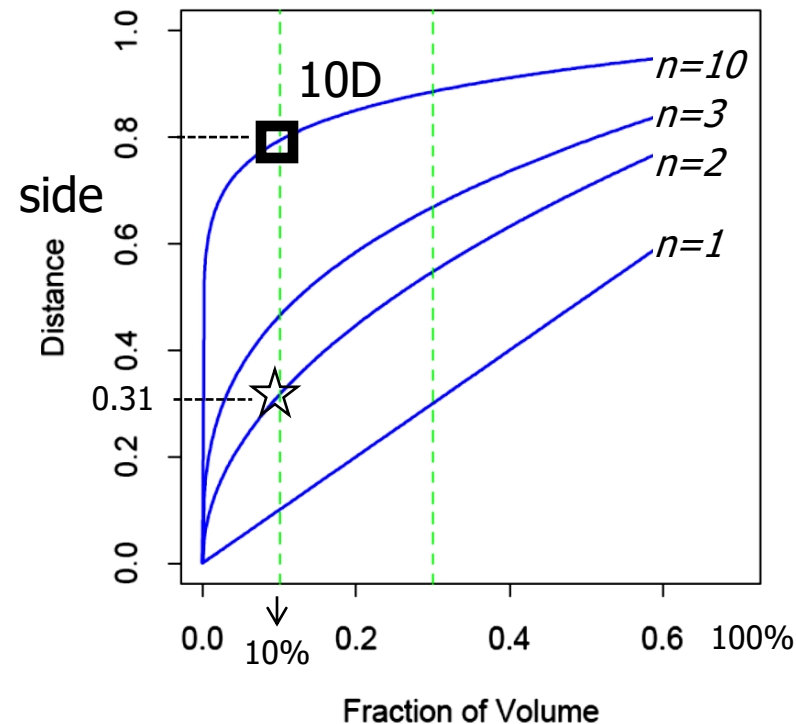



Figure 2.6: The curse of dimensionality is well illustrated by a subcubical neighborhood for uniform data in a unit cube.

The figure on the right shows the **sidelength** ( $=r^{1/n}$ ) of the subcube needed to capture a **fraction  $r$**  of the volume of the data, for different dimensions  $n$ .



In ten dimensions we need to cover 80%  $\blacksquare$  of the range of each coordinate to capture 10% of the data (since  $0.1^{1/10} \approx 0.8$ ), while in 2D 30% (0.316 sidelength) was suff.  $\star$

# Further explanation

## Loosing of generalization capability

1. Image that to have  $k$  data you need 10% of the volume. How much *sidelength* (features values range) do you need?
  - From 30% to 80% moving from 2D to 10D (**loosing similarity and locality!!!**)
  - And for 1% of the volume in 10D you still need 63% of the input range
2. On the other side:
  - 1D: with 0.3 *sidelength* we take  $\sim 30\%$  of data volume
  - 2D: with 0.3 *sidelength* we take 9% ( $\sim 10\%$ ) of data volume (the **red square**)
  - 3D: with 0.3 *sidelength* we take  $\sim 2.7\%$  of data volume (the **red cube**)
  - 10D: with 0.3 *sidelength* we take  $\sim 0.0006\%$  of data volume.

This sidelength can be not sufficient to get  $k$  data, unless we use **small K** (which, on the other hand, can lead to overfitting)



# (some) Limits of K-nn: curse of dim.

## 2. Low sampling density for high-dim data

- Sampling density ( $l/\text{volume}$ ) is proportional to  $l^{1/n}$  ( $l$  data,  $n$ =data dim). Hence, for example:
  - if 100 points are sufficient to estimate a function in  $IR^1$  (1-dim input),
  - $100^{10}$  ( $=100.000.000.000.000.000.000.000$ ) are needed to achieve similar accuracy in  $IR^{10}$  (10-dim input)

# (some) Limits of K-nn: curse of dim.

---



Dip. Informatica  
University of Pisa

## *3. Irrelevant features: The Curse of Noisy*

if the target depends on only few of many features in  $\mathbf{x}$  (e.g. 2 out 20), we could retrieve a “similar pattern” with the similarity dominated by the large number of irrelevant features

This issue grows with the dimensionality

# An improvement

---

Irrelevant features:

- We may weight features according to their relevance
  - Stretching the axes along some dimension
  - Weights can be searched by a (expensive) model selection approach or other approaches ...
- Feature selection approaches: it eliminates some variables → reduce input dimension

# Summary:

## K-nn design choices



Dip. Informatica  
University of Pisa

- The metric  $d$  to measure the closeness between patterns (e.g. Euclidian, Hamming, Manhattan distance..., weights on input features).  
Often this is the *key* for a successful application!
- K (number of neighbors: control underfitting/overfitting)

Often necessary to select:

- A subset of data (set of prototypes): e.g. by clustering
- A subset of features

Various approaches to deal with these issues.

# Extension in ML



Dip. Informatica  
University of Pisa

- Extensions to other local models in ML
  - Kernel smoothers
  - Local linear regression
  - Prototype methods
  - Case-based reasoning



# K-nn in the course: some general lessons



Dip. Informatica  
University of Pisa

- Too low variance is poor (rigid linear models), too high variance is dangerous (K-nn)
- Other concepts presented today that are interesting in general:
  - **Smoothing** techniques (in K-nn by increasing K)
  - **Curse of dimensionality** (the volume of the problem space increases so fast that the available data become sparse)
  - Again an **instance of the Statistical Learning Theory** bound plot: see the misclassification plot moving K
  - **Inductive Bias** for K-nn (metric based issue, which is underestimated in many books)

# ML Course structure

## And now?



Dip. Informatica  
University of Pisa

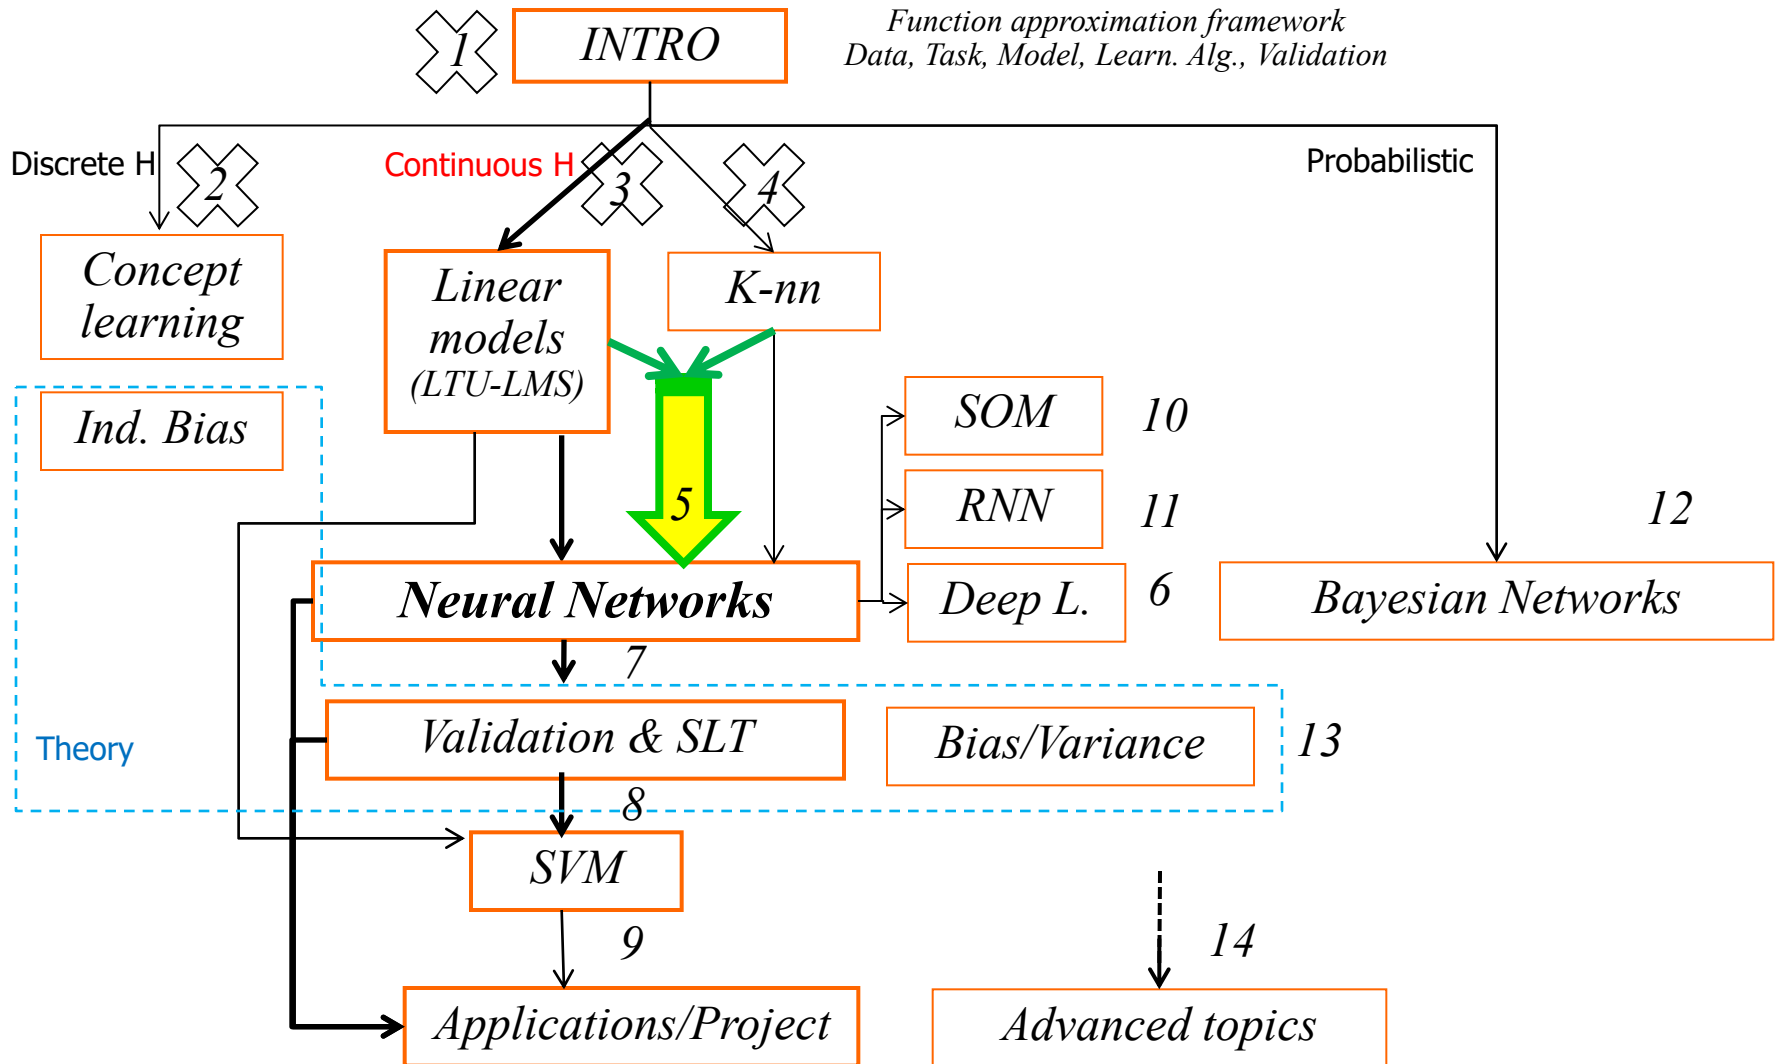
- K-nn does not build a “learned model”, not regularity/knowledge extraction/synthesis
  - It does not match the “learning” objective (that can forget the examples after the model is built)
- In the next lectures, we will look at model
  - Compact as the LTU model (all the knowledge in few parameters)
  - More flexible than the linear model (flexible as the K-nn)
    - with a suitable support to the control of the complexity

# ML Course structure

## Where we go



Dip. Informatica  
University of Pisa



# Bibliography: Linear and K-nn (I)



Dip. Informatica  
University of Pisa

- Hastie, Tibshirani, Friedman, The Elements of Statistical Learning, Springer Verlag, 2001-2017 (check the new Ed.): **chap 2**  
(note: exists an on-line pdf version)
- Mitchell:
  - Linear model and LMS alg.: chap 4.4
  - K-nn: **chap 8**
- Haykin (2nd edition):
  - Linear model and LMS alg.: chap 3 up to 3.7  
[details on Newton and Gauss-Newton methods are not strictly needed]
- Haykin (3rd edition):
  - Linear model and LMS alg.: chap 3  
[details on Newton and Gauss-Newton and other advanced approaches are not strictly needed]

# Bibliography: Linear and K-nn (II)



Dip. Informatica  
University of Pisa

- Moreover...
- Further readings (not mandatory!):
  - Gently introduction to linear models:  
AIMA (Russel, Norvig, Artificial Intelligence: A Modern Approach), ed .3: **chap 18.6** (18.6.1,18.6.2,18.6.3)

- To go in deep for **Linear least squares**

You can start from www resource as (With nice examples):

[https://en.wikipedia.org/wiki/Linear\\_least\\_squares\\_%28mathematics%29](https://en.wikipedia.org/wiki/Linear_least_squares_%28mathematics%29)

- **LS in general**

[http://en.wikipedia.org/wiki/Least\\_squares](http://en.wikipedia.org/wiki/Least_squares)

# Suggestion for the next lectures

---



Dip. Informatica  
University of Pisa

- Have a look to the slide of the next lectures and
- in particular to the proof of the Perceptron Convergence Theorem

# For information

**Alessio Micheli**  
**micheli@di.unipi.it**

[www.di.unipi.it/groups/ciml](http://www.di.unipi.it/groups/ciml)



Dipartimento di Informatica  
Università di Pisa - Italy  
Micheli



**Computational Intelligence &  
Machine Learning Group**